

Project Title: Stock Price Prediction for Netflix using LSTM

Problem Statement:

The stock market is a dynamic and complex system where investors aim to make informed decisions to maximize their returns. Predicting the stock price of a company like Netflix is particularly challenging due to its exposure to various market dynamics, competition, and global events. This project aims to develop a robust predictive model using Long Short-Term Memory (LSTM) neural networks to forecast the future stock prices of Netflix accurately.

Description:

Netflix is a leading streaming service provider, and its stock price is influenced by a myriad of factors, including subscriber growth, content releases, market sentiment, and economic conditions. This project will utilize historical stock price data, along with relevant features such as trading volume, macroeconomic indicators, and news sentiment analysis, to train an LSTM-based deep learning model.

Conclusion:

Summarize the project's findings and key results. Discuss the model's accuracy and its ability to predict Netflix stock prices. Highlight any challenges encountered during the project and potential areas for improvement. Emphasize the importance of using LSTM neural networks for time series forecasting tasks like stock price prediction.

Outline :

1. Libraries and settings
2. Analyze data
3. Manipulate data
4. Model and validate data
5. Predictions

```
# Disable all warnings
import warnings
warnings.filterwarnings('ignore')
import numpy as np    # linear algebra
import pandas as pd   # data processing, CSV file I/O (e.g.
pd.read_csv)
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
from sklearn.metrics import r2_score
import tensorflow
from tensorflow import keras
from keras.models import Sequential
from keras.layers import Dense
```

```

from keras.layers import SimpleRNN
from keras.layers import LSTM
from keras.layers import Dropout

# import all stock prices
data= pd.read_csv("/content/NFLX.csv", index_col = 0)
data.info()
data.head()

```

```

<class 'pandas.core.frame.DataFrame'>
Index: 1009 entries, 2018-02-05 to 2022-02-04
Data columns (total 6 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Open        1009 non-null   float64
1   High        1009 non-null   float64
2   Low         1009 non-null   float64
3   Close       1009 non-null   float64
4   Adj Close   1009 non-null   float64
5   Volume      1009 non-null   int64
dtypes: float64(5), int64(1)
memory usage: 55.2+ KB

```

	Open	High	Low	Close	Adj Close
2018-02-05	262.000000	267.899994	250.029999	254.259995	254.259995
2018-02-06	247.699997	266.700012	245.000000	265.720001	265.720001
2018-02-07	266.579987	272.450012	264.329987	264.559998	264.559998
2018-02-08	267.079987	267.619995	250.000000	250.100006	250.100006
2018-02-09	253.850006	255.800003	236.110001	249.470001	249.470001

	Volume
2018-02-05	11896100
2018-02-06	12595800
2018-02-07	8981500
2018-02-08	9306700
2018-02-09	16906900

```
data.tail()
```

	Open	High	Low	Close	Adj Close

Date

2022-01-31	401.970001	427.700012	398.200012	427.140015	427.140015
2022-02-01	432.959991	458.480011	425.540009	457.130005	457.130005
2022-02-02	448.250000	451.980011	426.480011	429.480011	429.480011
2022-02-03	421.440002	429.260010	404.279999	405.600006	405.600006
2022-02-04	407.309998	412.769989	396.640015	410.170013	410.170013

Volume

Date

2022-01-31	20047500
2022-02-01	22542300
2022-02-02	14346000
2022-02-03	9905200
2022-02-04	7782400

data.describe()

	Open	High	Low	Close	Adj Close
\					
count	1009.000000	1009.000000	1009.000000	1009.000000	1009.000000
mean	419.059673	425.320703	412.374044	419.000733	419.000733
std	108.537532	109.262960	107.555867	108.289999	108.289999
min	233.919998	250.649994	231.229996	233.880005	233.880005
25%	331.489990	336.299988	326.000000	331.619995	331.619995
50%	377.769989	383.010010	370.880005	378.670013	378.670013
75%	509.130005	515.630005	502.529999	509.079987	509.079987
max	692.349976	700.989990	686.090027	691.690002	691.690002

Volume

count	1.009000e+03
mean	7.570685e+06
std	5.465535e+06
min	1.144000e+06
25%	4.091900e+06
50%	5.934500e+06
75%	9.322400e+06
max	5.890430e+07

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 1009 entries, 2018-02-05 to 2022-02-04
Data columns (total 6 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Open        1009 non-null   float64
1   High        1009 non-null   float64
2   Low         1009 non-null   float64
3   Close       1009 non-null   float64
4   Adj Close   1009 non-null   float64
5   Volume      1009 non-null   int64
dtypes: float64(5), int64(1)
memory usage: 55.2+ KB
```

```
# Let's check the data, to see if there is duplicate data or not
data.duplicated()
```

```
Date
2018-02-05    False
2018-02-06    False
2018-02-07    False
2018-02-08    False
2018-02-09    False
...
2022-01-31    False
2022-02-01    False
2022-02-02    False
2022-02-03    False
2022-02-04    False
Length: 1009, dtype: bool
```

```
# Checking for missing values
data.isnull().sum()
```

```
Open      0
High      0
Low       0
Close     0
Adj Close 0
Volume    0
dtype: int64
```

Data Visualization

```
df = data.copy()
df
```

\	Open	High	Low	Close	Adj	Close
Date						
2018-02-05	262.000000	267.899994	250.029999	254.259995	254.259995	
2018-02-06	247.699997	266.700012	245.000000	265.720001	265.720001	
2018-02-07	266.579987	272.450012	264.329987	264.559998	264.559998	
2018-02-08	267.079987	267.619995	250.000000	250.100006	250.100006	
2018-02-09	253.850006	255.800003	236.110001	249.470001	249.470001	
...
2022-01-31	401.970001	427.700012	398.200012	427.140015	427.140015	
2022-02-01	432.959991	458.480011	425.540009	457.130005	457.130005	
2022-02-02	448.250000	451.980011	426.480011	429.480011	429.480011	
2022-02-03	421.440002	429.260010	404.279999	405.600006	405.600006	
2022-02-04	407.309998	412.769989	396.640015	410.170013	410.170013	
	Volume					
Date						
2018-02-05	11896100					
2018-02-06	12595800					
2018-02-07	8981500					
2018-02-08	9306700					
2018-02-09	16906900					
...	...					
2022-01-31	20047500					
2022-02-01	22542300					
2022-02-02	14346000					
2022-02-03	9905200					
2022-02-04	7782400					
[1009 rows x 6 columns]						
# Matrix form for correlation data						
drrr= data.corr()						
drrr						
	Open	High	Low	Close	Adj Close	Volume
Open	1.000000	0.998605	0.998508	0.996812	0.996812	-0.415838
High	0.998605	1.000000	0.998203	0.998551	0.998551	-0.400699
Low	0.998508	0.998203	1.000000	0.998544	0.998544	-0.432116

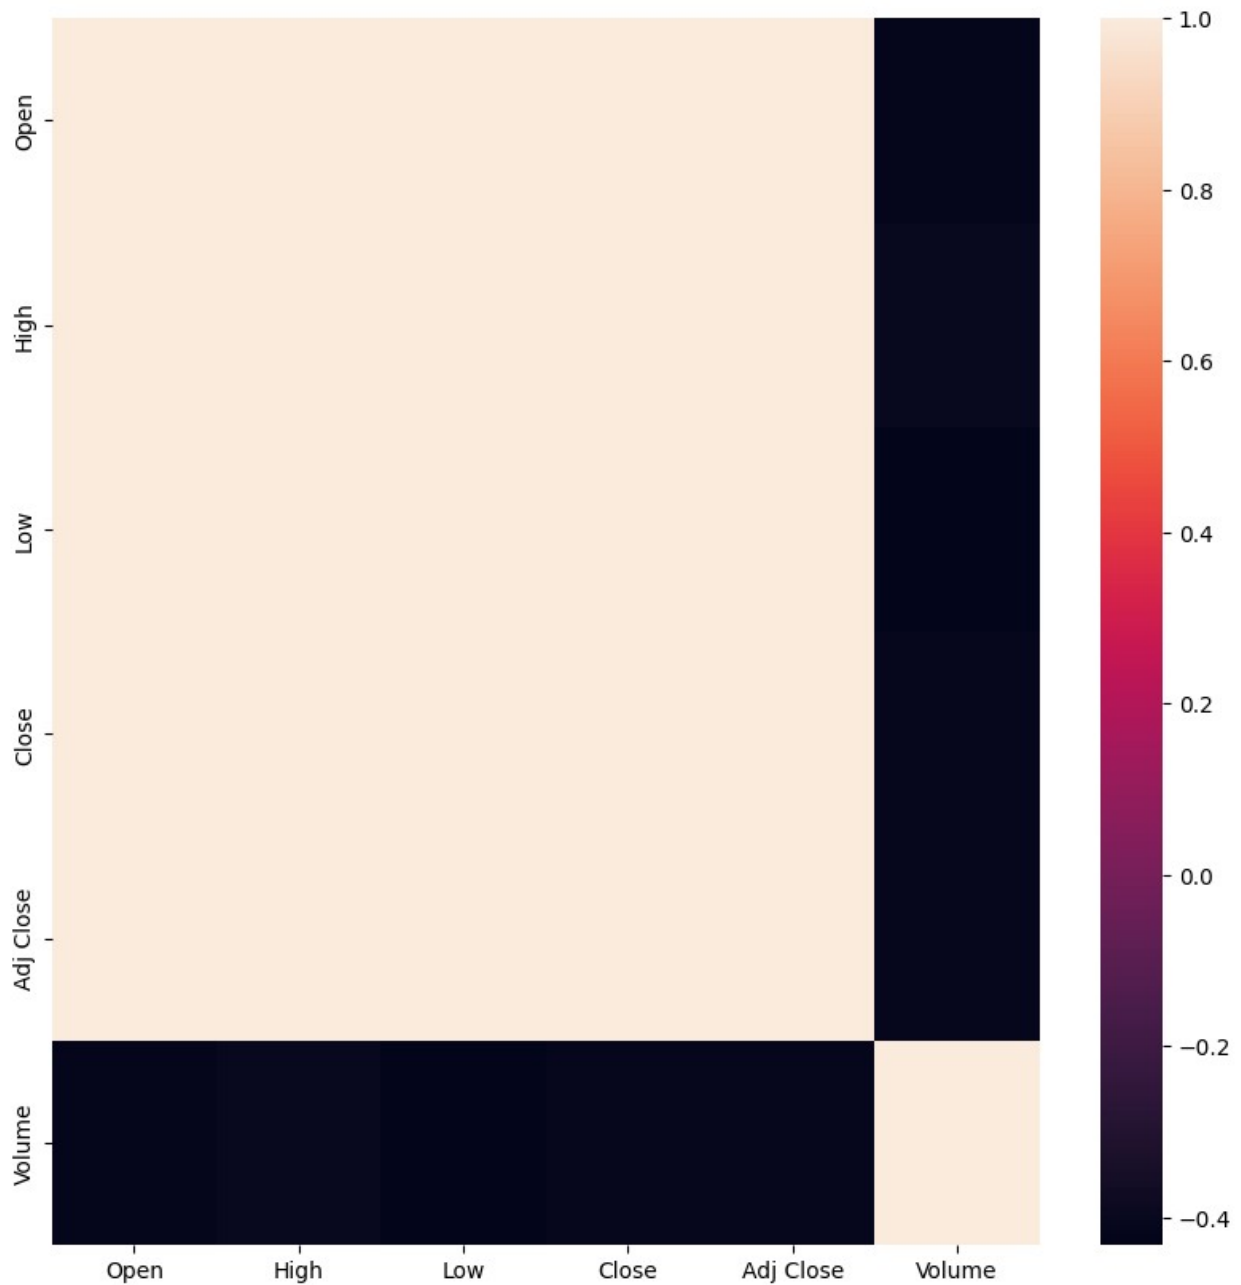
Close	0.996812	0.998551	0.998544	1.000000	1.000000	-0.413362
Adj Close	0.996812	0.998551	0.998544	1.000000	1.000000	-0.413362
Volume	-0.415838	-0.400699	-0.432116	-0.413362	-0.413362	1.000000

We here looking at the data Visualization by heatmap.

```
plt.figure(figsize=(10,10))
```


```
sns.heatmap(drrr)
```

<Axes: >




```
f,ax = plt.subplots(figsize=(10, 10))
sns.heatmap(drrr, annot=True, linewidths=.5, fmt= '.1f',ax=ax)
```


<Axes: >




```
# Original data.
for i in df.columns[~df.columns.isin(["Date", "Volume"])]:
    f = sns.countplot(x=df[i]);
    plt.figure(figsize=(10,10))
    plt.show()
```




<Figure size 1000x1000 with 0 Axes>




<Figure size 1000x1000 with 0 Axes>



<Figure size 1000x1000 with 0 Axes>




<Figure size 1000x1000 with 0 Axes>



<Figure size 1000x1000 with 0 Axes>

```
df.plot(kind='box',subplots=True,layout=(10,10),figsize=(20,20))
plt.show()
```



Splitting

```
train_data, test_data = train_test_split(data, test_size=0.2,
random_state=42)
train_data.shape, test_data.shape
((807, 6), (202, 6))
```

Preprocessing

```

scaler = MinMaxScaler()
scalerr = MinMaxScaler()
train_data_scaled = scaler.fit_transform(train_data)
train_data_scaled

array([[0.17860959, 0.14960206, 0.16969522, 0.17439697, 0.17439697,
        0.14992077],
       [0.28898633, 0.26224888, 0.27997145, 0.27364546, 0.27364546,
        0.10854764],
       [0.30501932, 0.28738923, 0.30877797, 0.30631203, 0.30631203,
        0.06049766],
       ...,
       [0.67691906, 0.66287127, 0.6698401 , 0.66183209, 0.66183209,
        0.02491101],
       [0.09626332, 0.07914141, 0.10282939, 0.10512013, 0.10512013,
        0.08610638],
       [0.33054124, 0.33309287, 0.3316983 , 0.35996755, 0.35996755,
        0.1189803 ]])

```

```

test_data_scaled = scaler.transform(test_data)
test_data_scaled

```

```

array([[0.57747095, 0.5857591 , 0.59624089, 0.60416715, 0.60416715,
        0.07499163],
       [0.59230423, 0.58091143, 0.56955252, 0.57150058, 0.57150058,
        0.08023833],
       [0.5964488 , 0.59491335, 0.59831441, 0.58493084, 0.58493084,
        0.05016219],
       ...,
       [0.11825144, 0.10166625, 0.09696552, 0.09081345, 0.09081345,
        0.19982609],
       [0.28200602, 0.27101986, 0.29424097, 0.29827134, 0.29827134,
        0.04948878],
       [0.2564841 , 0.23776243, 0.25747474, 0.25381769, 0.25381769,
        0.14632288]])

```

#Creating a data structure with 50 timesteps and 1 output, timestep is our memory size

#In this function we are creating our train data with 50x stock price and next one is 1 scrolled data.

#for example X_train[0] will be our data's 0 to 49. values

#X_train[1] will be our data's 1 to 50. values

#this 50 is our memory size, it will remember this way what we had before.

```
X_train=[]
```

```
y_train=[]
```

```
timesteps=10
```

```
for i in range(timesteps,len(train_data_scaled)):
```

```
    X_train.append(train_data_scaled[i-timesteps:i,0])
```

```
    y_train.append(train_data_scaled[i,0])
```

```
X_train,y_train=np.array(X_train),np.array(y_train)
real_stock_price=test_data.loc[:,["Open"]].values
```

Model

Recurrent Neural Network (RNN)

```
regressor=Sequential()
#Adding the first RNN Layer and some Dropout Regularisation
regressor.add(SimpleRNN(units=10,activation="tanh",return_sequences=True,input_shape=(X_train.shape[1],1)))
regressor.add(Dropout(0,2))

#Adding the second RNN Layer and some Dropout Regularisation
regressor.add(SimpleRNN(units=10,activation="tanh",return_sequences=True))
regressor.add(Dropout(0,2))

#Adding the third RNN Layer and some Dropout Regularisation
regressor.add(SimpleRNN(units=10,activation="tanh",return_sequences=True))
regressor.add(Dropout(0,2))

#Adding the third RNN Layer and some Dropout Regularisation
regressor.add(SimpleRNN(units=10,activation="tanh",return_sequences=True))
regressor.add(Dropout(0,2))

#Adding the fourth RNN Layer and some Dropout Regularisation
regressor.add(SimpleRNN(units=10))
regressor.add(Dropout(0,2))

#Adding the output Layer
regressor.add(Dense(units=1))

regressor.summary()
```

Model: "sequential_4"

Layer (type)	Output Shape	Param #
=====		
simple_rnn_15 (SimpleRNN)	(None, 50, 10)	120
dropout_21 (Dropout)	(None, 50, 10)	0
simple_rnn_16 (SimpleRNN)	(None, 50, 10)	210
dropout_22 (Dropout)	(None, 50, 10)	0

simple_rnn_17 (SimpleRNN)	(None, 50, 10)	210
dropout_23 (Dropout)	(None, 50, 10)	0
simple_rnn_18 (SimpleRNN)	(None, 50, 10)	210
dropout_24 (Dropout)	(None, 50, 10)	0
simple_rnn_19 (SimpleRNN)	(None, 10)	210
dropout_25 (Dropout)	(None, 10)	0
dense_4 (Dense)	(None, 1)	11

=====

Total params: 971 (3.79 KB)
Trainable params: 971 (3.79 KB)
Non-trainable params: 0 (0.00 Byte)

```
#Compiling the RNN
regressor.compile(optimizer="adam",loss="mean_squared_error")
#Fitting the RNN to the Training set
regressor.fit(X_train,y_train,epochs=10,batch_size=32)
```

```
Epoch 1/10
24/24 [=====] - 5s 45ms/step - loss: 0.1493
Epoch 2/10
24/24 [=====] - 2s 67ms/step - loss: 0.0667
Epoch 3/10
24/24 [=====] - 2s 79ms/step - loss: 0.0613
Epoch 4/10
24/24 [=====] - 1s 45ms/step - loss: 0.0593
Epoch 5/10
24/24 [=====] - 1s 44ms/step - loss: 0.0589
Epoch 6/10
24/24 [=====] - 1s 45ms/step - loss: 0.0580
Epoch 7/10
24/24 [=====] - 1s 45ms/step - loss: 0.0591
Epoch 8/10
24/24 [=====] - 1s 44ms/step - loss: 0.0575
Epoch 9/10
24/24 [=====] - 1s 45ms/step - loss: 0.0569
Epoch 10/10
24/24 [=====] - 1s 45ms/step - loss: 0.0567
```

<keras.src.callbacks.History at 0x7f37d12f4eb0>

Long Short Term Memory (LSTM)

```

model = Sequential()
model.add(LSTM(units=50, return_sequences=True,
input_shape=(X_train.shape[1], 1)))
model.add(Dropout(0.2))
model.add(LSTM(units=50, return_sequences=True))
model.add(Dropout(0.2))
model.add(LSTM(units=50, return_sequences=True))
model.add(Dropout(0.2))
model.add(LSTM(units=50, return_sequences=True))
model.add(Dropout(0.2))
model.add(LSTM(units=50, return_sequences=True))
model.add(Dropout(0.2))
model.add(LSTM(units=50))
model.add(Dropout(0.2))
model.add(Dense(units = 1))

```

```
model.summary()
```

Model: "sequential_5"

Layer (type)	Output Shape	Param #
lstm_6 (LSTM)	(None, 50, 50)	10400
dropout_26 (Dropout)	(None, 50, 50)	0
lstm_7 (LSTM)	(None, 50, 50)	20200
dropout_27 (Dropout)	(None, 50, 50)	0
lstm_8 (LSTM)	(None, 50, 50)	20200
dropout_28 (Dropout)	(None, 50, 50)	0
lstm_9 (LSTM)	(None, 50, 50)	20200
dropout_29 (Dropout)	(None, 50, 50)	0
lstm_10 (LSTM)	(None, 50, 50)	20200
dropout_30 (Dropout)	(None, 50, 50)	0
lstm_11 (LSTM)	(None, 50)	20200
dropout_31 (Dropout)	(None, 50)	0
dense_5 (Dense)	(None, 1)	51

Total params: 111451 (435.36 KB)
Trainable params: 111451 (435.36 KB)
Non-trainable params: 0 (0.00 Byte)

```
model.compile(loss = 'mean_squared_error', optimizer = 'Adam')  
model.fit(X_train, y_train, epochs = 10, batch_size = 32)
```

```
Epoch 1/10  
24/24 [=====] - 20s 149ms/step - loss: 0.0890  
Epoch 2/10  
24/24 [=====] - 4s 187ms/step - loss: 0.0603  
Epoch 3/10  
24/24 [=====] - 4s 160ms/step - loss: 0.0607  
Epoch 4/10  
24/24 [=====] - 4s 152ms/step - loss: 0.0601  
Epoch 5/10  
24/24 [=====] - 4s 156ms/step - loss: 0.0598  
Epoch 6/10  
24/24 [=====] - 5s 188ms/step - loss: 0.0588  
Epoch 7/10  
24/24 [=====] - 4s 148ms/step - loss: 0.0593  
Epoch 8/10  
24/24 [=====] - 4s 151ms/step - loss: 0.0587  
Epoch 9/10  
24/24 [=====] - 5s 202ms/step - loss: 0.0590  
Epoch 10/10  
24/24 [=====] - 4s 151ms/step - loss: 0.0587
```

<keras.src.callbacks.History at 0x7f37cf5c2980>

```
dataset_total=pd.concat((train_data["Open"],test_data["Open"]),  
axis=0)  
inputs= dataset_total[len(dataset_total)-len(train_data)-  
timesteps:].values.reshape(-1,1)  
inputs=scalerr.fit_transform(inputs)
```

#prediction

```
X_test=[]  
for i in range(timesteps,timesteps+len(test_data)):  
    X_test.append(inputs[i-timesteps:i,0])  
X_test=np.array(X_test)  
X_test=np.reshape(X_test,(X_test.shape[0],X_test.shape[1],1))  
predicted_stock_price=model.predict(X_test)  
predicted_stock_price=scalerr.inverse_transform(predicted_stock_price)  
#we had scaled between 0-1 data, inversing it
```

```
7/7 [=====] - 4s 40ms/step
```

#visualising

```
plt.plot(real_stock_price,color="red",label="Real Netflix Stock
```

```
Price")
plt.plot(predicted_stock_price,color="blue",label="Predicted Netflix
Stock Price")
plt.title("Netflix Stock Price Prediction")
plt.xlabel("Time")
plt.ylabel("Stock Price")
plt.legend()
plt.show()
```

